

Vrije Universiteit Amsterdam



Bachelor Thesis

---

# RDMA Key Value store: Scalability of RDMA transportation types

---

**Author:** Jan Erik (Jens) Troost (2656054)

*1st supervisor:* Animesh Trivedi  
*daily supervisor:* supervisor name  
*2nd reader:* supervisor name

*A thesis submitted in fulfillment of the requirements for  
the VU Bachelor of Science degree in Computer Science*

May 29, 2021

## Abstract

Here goes the abstract of this thesis.

# Contents

<b>List of Figures</b>	<b>iii</b>
<b>List of Tables</b>	<b>v</b>
<b>1 Introduction</b>	<b>1</b>
1.1 Context . . . . .	1
1.2 Problem statement . . . . .	2
1.3 Research Question . . . . .	2
1.4 Research Method . . . . .	2
1.5 Plagiarism Declaration . . . . .	3
1.6 Thesis Structure . . . . .	3
<b>2 Background</b>	<b>5</b>
2.1 rdma . . . . .	5
<b>3 Design</b>	<b>7</b>
<b>4 Implementation</b>	<b>9</b>
<b>5 Evaluation</b>	<b>11</b>
<b>6 Related Work</b>	<b>13</b>
<b>7 Conclusion</b>	<b>15</b>
<b>References</b>	<b>17</b>

## CONTENTS

---

# List of Figures

## LIST OF FIGURES

---

# List of Tables

## LIST OF TABLES

---



# 1

## Introduction

### 1.1 Context

Seemingly ever growing tech giants, such as Facebook, Amazon and Google, require fast, reliable and scalable key-value storage (KV-store) to serve product recommendations, user preferences, and advertisements. A KV-store is a data storage type which aims to store many small values and provides a high throughput and low latency interface. KV-stores usually offer a simple interface using 3 commands: GET, SET, and DEL. A trivial implementation makes use of a hash-table which matches the hash value of a given key, to the value associated with that key. Past advancements of KV-stores have focussed on improving these underlying data structures. For example, Cuckoo and Hopscotch hashing(2), have presented an alternative of an improved hash table algorithm. Nonetheless, with recent developments with RDMA networks, this has become a more attractive direction to improve KV-stores.

Remote direct memory access (RDMA) networks have increasingly become more popular in commercial and academic data centers, due to the increase in availability and decrease in cost for RDMA capable network interface cards (RNICs). RDMA offers a more direct connection between two machines, by interfacing the memory directly.

Instead of the operating system (OS) handling the incoming and outgoing packets, RNICs implement this in hardware, requiring no additional help from the OS and CPU. This results in lowering the overall latency and CPU overhead.

## 1. INTRODUCTION

---

### 1.2 Problem statement

Making efficient use of RDMA networks is a difficult task, and requires in-depth knowledge on the hardware constraints present in the RNICs (1). Additionally, in a higher level sense, there are design choices to be made. RDMA networks can handle of various transportation types and so-called verbs. Each of these have advantage and disadvantages, in general setting, which can be read about in 2.1.

Furthermore, with the continuous growth of tech giants, an important factor is scalability. Some of these design choices impact scalability.

### 1.3 Research Question

This paper will explore to what extent RDMA transportation types affect the performance and scalability of KV-store. For this, this research will answer the following questions:

RQ1: What are the advantages and disadvantages of RDMA transportation types on an KV-store? This question aims to apply the already known information on the RDMA transportation types, on a RDMA KV-store specific setting. Aiding the design processes of further RDMA KV-store implementation, and possibly related implementations.

RQ2: How scalable are these transportation types? This being an important factor for tech giants which make use of RDMA KV-store, and require a system to be scalable and reach multiple clients.

### 1.4 Research Method

M1: A KV-store prototype will be implemented, with a flexible network interfacing, to include both socket and RDMA interfacing. This implementation can not only be used for this paper, but could also be expanded upon, for example with more advanced KV-store implementations.

M2: To investigate the performance and scalability between the various networking types, macro level benchmark will be designed, with realistic KV-store workloads.

A simple KV-store will be implemented. The backend, networking side, will be flexibly implemented to accommodate for the various transportation types.

Benchmarks will be done on DAS 5 computing cluster.

## 1.5 Plagiarism Declaration

## 1.6 Thesis Structure

Section 2 will provide the necessarily background knowledge of KV-store, linux sockets, and RDMA. Next, section 3 will go over the design of the KV-store, networking interfacing, and benchmark. In section 4 this design will be taken and described the implementation in a more technical prospective. The benchmarking result will be presented and analyzed in section 5. Section 6 compares findings with that of previously done work. Closing off the thesis, section 7 will go over the conclusion and provide recommendations.

## 1. INTRODUCTION

---

## 2

# Background

...

### 2.1 RDMA

## 2. BACKGROUND

---

3

Design

...

### 3. DESIGN

---



4

# Implementation

...

#### 4. IMPLEMENTATION

---

5

## Evaluation

...

## 5. EVALUATION

---

6

## Related Work

...

## 6. RELATED WORK

---

7

## Conclusion

...

## 7. CONCLUSION

---



# References

- [1] Youmin Chen, Youyou Lu, and Jiwu Shu. Scalable rdma rpc on reliable connection with efficient resource sharing. In *Proceedings of the Fourteenth EuroSys Conference 2019*, pages 1–14, 2019. 2
- [2] Roxana Geambasu, Amit A Levy, Tadayoshi Kohno, Arvind Krishnamurthy, and Henry M Levy. Comet: An active distributed key-value store. In *OSDI*, pages 323–336, 2010. 1