

Vrije Universiteit Amsterdam



Bachelor Thesis

Scalability of RDMA transportation types in a Key Value store application

Author: Jan Erik (Jens) Troost (2656054)

1st supervisor: Dr. Animesh Trivedi
daily supervisor: Dr. Animesh Trivedi
2nd reader: Jesse Donkervliet

*A thesis submitted in fulfillment of the requirements for
the VU Bachelor of Science degree in Computer Science*

June 25, 2021

Abstract

Seemingly ever growing tech giants, such as Facebook, Amazon and Google, require fast, reliable and scalable key-value storage (KV-store) to serve product recommendations, user preferences, and advertisements. Past advancements of KV-stores have focussed on improving these underlying data structures. Remote direct memory access (RDMA) networks have increasingly become more popular in commercial and academic data centers. This relatively new technology offers lower latency and higher throughput performance compared to traditional sockets. Scalability performance of RDMA transportation protocols has seen relatively little investigation, the majority of previous work focused on RDMA verb choice.

This paper focuses on this gap, evaluating scalability performance of RDMA transportation types. These findings will be used to give recommendation for RDMA KV-store implementations. Macro-level benchmarks have been conducted, with realistic KV-store workloads. This has shown that RDMA offers a significant improvements compared to TCP. Unreliable datagram (UD) has shown to perform best, with 94.8% throughput improvement over TCP, and 58.7% improvement over reliable connection (RC). All the while offering consistent and gradually increasing latency up to 30 clients, reaching a maximum average latency of $64 \mu\text{sec}$, which is an 52.0% improvement over TCP. Additionally, scalability issues of RC have been shown, and is not recommended in a KV-store application.

Contents

List of Figures	iii
List of Tables	v
1 Introduction	1
1.1 Context	1
1.2 Problem statement	2
1.3 Research Question	2
1.4 Research Methods	3
1.5 Thesis Contributions	3
1.6 Plagiarism Declaration	3
1.7 Thesis Structure	4
2 Background	5
2.1 KV-store	5
2.2 Linux socket and TCP	5
2.3 RDMA	6
2.3.1 Queue Pairs	6
2.3.2 Transportation types	7
2.3.3 Verbs	8
2.3.4 Connecting Queue Pairs	8
2.3.5 Programming with RDMA	9
3 Designing a RDMA Key Value store	11
3.1 Key value store	11
3.1.1 Requests	11
3.1.2 Response	11
3.1.3 Hash table	13

CONTENTS

3.1.4	Multithreaded	13
3.2	RDMA	14
3.2.1	Queue pairs	14
3.3	Benchmark design	14
3.3.1	Experimental Setup	15
4	Evaluation	17
4.1	Throughput analysis	18
4.1.1	Scalability rate	19
4.2	Latency analysis	19
4.2.1	Variation from average	21
4.2.2	Distribution of average latency	23
4.2.3	All 30 clients and per operation latency	23
4.2.4	RDMA multiple QP's causing cache misses and context switching	23
4.2.4.1	Cache misses as increased variation for RC	23
5	Related Work	27
5.1	HERD	27
6	Future improvements	29
6.1	Experimental transport types	29
6.2	Optimizations	29
6.3	Newer hardware	30
7	Conclusion	31
References		35

List of Figures

2.1	Process overview of <i>SEND</i> and <i>RECV</i>	9
3.1	Structure for request with accommodated method	12
3.2	Structure for response with accommodated response codes	12
3.3	Key value server layout	13
4.1	Throughput of clients executing 10 million operations	18
4.2	Latency of clients executing 10 million operations	20
4.3	Structure for response with accommodated response codes	24

LIST OF FIGURES

List of Tables

2.1	Verbs available to each transportation type	8
3.1	DAS-5 cluster specifications	15
4.1	Throughput increase per additional client	19
4.2	Numeric statistic latency. All statistical values have unit μ sec.	22

LIST OF TABLES

1

Introduction

1.1 Context

Seemingly ever growing tech giants, such as Facebook, Amazon and Google, require fast, reliable and scalable key-value storage (KV-store) to serve a variety of services: product recommendations, user preferences, and advertisements[12, 16]. Amazon’s Dynamo is an example of such a highly used KV-store, which is used by Amazon to provide its core services, such as session storage and shopping cart[12]. At peak, these services handle tens of millions customers a second. This requires Dynamo to process requests in the span of a few milliseconds, to provide a consistent and fast user experience.

Remote direct memory access (RDMA) networks have increasingly become more popular in data centers, due to their increase in throughput, decrease in latency, and lower cost of RDMA capable network interface cards (RNIC)[11, 23]. RDMA and RNIC provides a lower latency networking interface compared to classical network cards, which is advantage for latency sensitive applications, such as a KV-store, deployed in data centers.

In a more general sense, a KV-store aims to store many small values, identified with a key. KV-stores usually offer a simple interface using 3 commands: *GET*, *SET*, and *DEL*. With these commands, data can be retrieved, updated/added, or deleted. Typically, the basis of a KV-store is formed by a hash-table, which pairs a unique key with a value. Past advancements of KV-stores have focussed on improving these underlying data structures[14, 25]. Cuckoo and Hopscotch hashing[13, 22] have presented an alternative of an improved hash table algorithm. Nonetheless, with recent developments with RDMA networks, this has become a more attractive direction to improve KV-stores[11].

Remote direct memory access (RDMA) networks have increasingly become more popular in commercial and academic data centers, due to the increase in availability and decrease

1. INTRODUCTION

in cost for RDMA capable network interface cards (RNICs)[11, 23]. RDMA offers a more direct connection between two machines, by using direct memory access (DMA), without the involvement of host CPU or OS.

1.2 Problem statement

Making efficient use of RDMA networks is a difficult task, and requires in-depth knowledge on the hardware constraints present in the RNICs[11, 23]. Additionally, in a higher level sense, there are design choices to be made, in regard to transportation types and so-called verbs. RDMA networks can handle of various transportation types: reliable connection (RC), unreliable connection (UC) and unreliable datagrams (UD). Additionally, verbs can be seen as network operations. There are four main verbs: *SEND*, *RECV*, *READ*, and *WRITE*. Each transportation type and verb come with advantages and disadvantages, which will be discussed further in this thesis. However, impact of transportation types on scalability of an RDMA bases KV-store has not been explicitly examined. Previous work has focused on researching verb choices, less on transportation type[22, 24, 27].

Scalability of their KV-store service is an important factor for tech giants such as Amazon. With their growth in number of customers and concurrent clients, their KV-store servers are hit with large number of clients.

Scalability is an important factor for growing tech giants, as their increasing customer . These require a low latency and scale to reach large number of clients, all the while being highly available[12], which RDMA could provide.

1.3 Research Question

This paper will explore to what extent RDMA transportation types affect the performance and multicore scalability of KV-store. For this, this research will answer the following questions:

RQ1 What is the multicore scalability of the RDMA transportation types: reliable connection (RC), unreliable connection (UC), and unreliable datagram (UD)? With the growing number of clients that

RQ2 What are the advantages and disadvantages of RDMA transportation types on an KV-store? This question aims to aid with design process of future RDMA KV stores. By applying known ramifications of RDMA to that in a KV store use case. With

supporting results, recommendations can be given, and possible design issues can be foreseen.

1.4 Research Methods

In this thesis the network performance of KV-store, with varying network protocols, will be studied. First an understanding of KV-stores is established. Along with discussing known issues with traditional networking, RDMA will be introduced. Experiments will be performed to measure overall throughput and latency, with increasing number of clients. With this the possible scalability of each transport type will be shown.

A prototype[17, 20, 30] and experimental[19, 21, 28] approach is taken for this thesis:

M1 A KV-store prototype will be implemented, with a flexible network interfacing, to include both traditional socket and RDMA interfacing. This, and all other relevant project files, are open source and can be found on Github[9].

M2 To investigate the performance and scalability between the various networking types, a workload realistic[10] macro-level benchmark will be designed.

With the focus of this thesis being on the network implementation, a trivial KV-store will be used. This implementation does not offer strong performance and scalability, however these issues are minimized and kept consistent throughout experiments.

All measurements hereafter are recorded on the DAS-5 computing cluster. This allows for a consistent working environment. Further details on this is shown in table 3.1 and discussed in section 3.3.1

1.5 Thesis Contributions

This thesis presents scalability performance of RC and UD are shown and compared to the well established TCP protocol. Additionally, recommendation as to which transportation protocol is best suited for an RDMA KV-store application. These findings can be used to further examine possible optimizations or RDMA verb choices, or aid in future RDMA KV-store implementations. The prototype and all project files can be found on Github[9].

1.6 Plagiarism Declaration

I hereby declare that this thesis work, all data collected and findings are my own.

1. INTRODUCTION

For this thesis, the basic structure of KV store from course Operating Systems, at the VU, has been used. Along with basic RDMA functions from Dr. Trivedi's RDMA example[7].

Note: permission incoming for use of KV store structure from OS course.

1.7 Thesis Structure

Section 2 will provide the necessarily background knowledge of KV-store, linux sockets, and RDMA. Next, section 3 will go over the design of the KV-store, networking interfacing, and benchmark. In section 3 this design will be taken and described the implementation in a more technical prospective. The benchmarking result will be presented and analyzed in section 4. Section 5 compares findings with that of previously done work. Future prospects will be given in section 6. Closing off the thesis, section 7 will go over the conclusion and provide recommendations.

2

Background

In this section, background information on key value stores, which is the backbone of this thesis, will be given. Also, the issues with commonly used Linux socket, will be explained. Further, RDMA is thoroughly explained, as this is crucial to the understanding of this paper, and how it addresses issues facing sockets.

2.1 Key Value Store

Key value stores are extensively used to offer low latency look ups. These have a wide variety of use cases, most notably as cache systems, such as Memcached[1], and as remote DRAM storage, such as RAMCloud[29]. In its simplest form, KV store use a set of commands, commonly *SET*, *GET*, and *DEL*, to perform tasks on a server. These commands interact with a fast data structure, like hash table or similar. Typically, this has been the target for advancements in KV stores, using lower latency, memory efficient, and scalable data structures[14, 25].

It has been found that typical work loads consist mostly of *GET* requests. Atikoglu et. al. analysed workloads on Memcached systems, and have found that on average a 30:1 (95%) *GET/SET* ratio[10]. This figure will be used in the benchmarking design, see section 3.3.

2.2 Linux sockets and TCP network stack

Linux sockets are versatile programming API, offering a simple interface to network communication. Behind this socket interface, the kernel is tasked with data and memory, and connection management[18, 32]. This results in CPU cycles being used to process

2. BACKGROUND

incoming packets and data copies, while these cycles could be used for other tasks. Linux takes an extensive route when dealing with packets, as shown by Hanford et. al. detailed investigation[18]. Memory copies have been found to cause delay in Frey and Alonso’s work[15], however for small message sizes, Yasukata, and team, found this is insignificant[34]. TCP processing accounted for more than half of the overhead, after subtracting unavoidable PCIe latencies.

Past attempts at improving TCP performance included offloading most processing to the network interface card (NIC)[18]. Higher end NICs use an TCP offload engine (TOE) and DMA. The socket API still requires system calls, among other limitations, thus requiring kernel trapping for management.

2.3 RDMA

Remote Direct Memory Access (RDMA) address the issues of linux sockets, providing lower latency, less CPU overhead, and potentially higher throughput. RDMA has been used in super computers for many years, however recently have seen significant improvements. RDMA capable network cards (RNICs) have seen lowering in cost [23], which made this an appealing improvement to data center networking.

RDMA achieves this low latency by offloading the network processing onto the RNIC, bypassing CPU and kernel entirely. As shown in section 2.2 above, the linux kernel has an extensive route, from application to NIC to network, including system calls and memory copies. With RDMA, a zero-copy memory access can be realised through DMA and programmable IO (PIO) operations, which can be seen in 2.1. This makes RDMA a compelling technology for latency sensitive workloads, like KV-store, making remote memory operations possible and perform near to local memory operation speeds.

2.3.1 Queue Pairs

RDMA is largely based around the notion of a queue pair (QP). These consist of a send and receive queue, these are the essence of performing network operations. These queues can be seen as the receive and transmit queue (RX and TX respectively) in classical NIC, however are bound per QP instead of being shared system wide. A QP is similar to a linux socket, as these are used to send and receive data. A work request (WR) is a type of command that tells the RNIC what task to perform and which memory locations are needed, these are passed to the RNIC via PIO operations. Every WR placed either in the send or receive queue will be consumed in the same order as placed in the queue, similarly

for completions of WR's. This is important when dealing with unreliable transportation, as will be explained in 2.3.2.

Queue pairs are also linked with a completion queue (CQ). This queue is used as a notification queue, with the status of WR's. For signalled WR's, a completion event will be passed from RNIC to CPU, thus minor CPU involvement.

For RDMA operations like *READ* and *WRITE* this would include the remote address to be accessed, with two-sided verbs such as *SEND* and *RECV* this is accompanied by a buffer for which DMA operations can take place. In this paper, the focus will be on *SEND* and *RECV* verbs, more can be read in section 2.3.3 below and in 3.2.

2.3.2 Transportation types

Much like traditional networking, RDMA can make use of several transport protocols. Simply put, these can be split into two types, connection and datagram, and reliable and unreliable. There are 3 main transportation types: reliable connection (RC), unreliable connection (UC), and unreliable datagram (UD).

Firstly, unreliable protocols do not make use of ACK/NAK packets, while this is used for reliable, this could result in packet loss and unordered packets. However, it has been shown that this rarely occurs[22, 24]. Additionally, this could require application retransmission handling. With reliable protocol this process is done by the RNIC, without OS or application involvement.

As stated in the name, RC and UC need a connection between queue pairs (QP). Only one QP can be connected to another QP. Contrasting this, unreliable datagram (UD) do not require a connection. This meaning that a UD QP can communicate with any other UD QP. UD therefore can make efficient use of a one-to-many network topology or application, and RNIC cache. Every QP has to be held by RNIC in cache, which is severely limited[31]. A cache miss or many context switching in RNIC, would diminish performance. With the fan-in and fan-out effect present in many data center workloads[33], the potential scalability of connection oriented QP's is limited[24]. UD requires an address handler (AH) to be passed along with its WR's. This is the route and the destination for a packet. The destination does not need to be a QP, but could also be a multicast group. Multiple clients can listen on the same channel, this can be applied in applications where the same data can be used by multiple clients. An AH can be created from a previous work completion (WC), or via metadata exchange, see 2.3.4.

UD is limited however in the maximum transmission unit (MTU). As can be seen in table 2.1, UD has a maximum message size of 4KB, beyond this and the message is divided into

2. BACKGROUND

	SEND	RECV	READ	WRITE	MTU size
<i>RC</i>	YES	YES	YES	YES	2 GB
<i>UC</i>	YES	YES	NO	YES	2 GB
<i>UD</i>	YES	YES	NO	NO	4 KB

Table 2.1: Verbs available to each transportation type

several packets. This strongly contrasts the 2 GB available for RC and UC. This has to be taken into account on a per-application bases. In the case of this thesis, 4 KB is sufficient message size for KV-stores. Frey and Alonso have comprised a list of application and if it is suitable for RDMA[15].

2.3.3 Verbs

To interact with RNICs, RDMA uses so-called "verbs" to execute specific types of instructions. Some of which are: read, write, send, and receive. Read and write (*READ* and *WRITE*) follow so-called memory semantics, while send and receive (*SEND* and *RECV*) follow channel semantics. Memory semantics require the destination memory address to be known. This meaning, to be able to perform a RDMA read of a remote memory location, the memory address of the requested memory needs to be known.

Channel semantics are similar to socket, in the sense that the remote memory address does not need to be known. However, to perform a *SEND* operation, the receiving end must post a *RECV* WR before the *SEND* WR is sent, this is called pre-posting. This tells the server's RNIC which memory location the application expects the next incoming message to be place.

Not all verbs are available to every QP type. Table 2.1 summarizes the transportation type and which verb is available.

2.3.4 Connecting Queue Pairs

Establishing a connection between QP's involves exchanging metadata related to the QP. This can be done via a known QP, or via traditional networks, using for example TCP. However, in some cases, a classical ethernet NIC is not available. In these circumstances, a known QP, Communication Manager (CM), can be used. This QP which is always available when using RNIC, and can be used to communicate between RNICs.

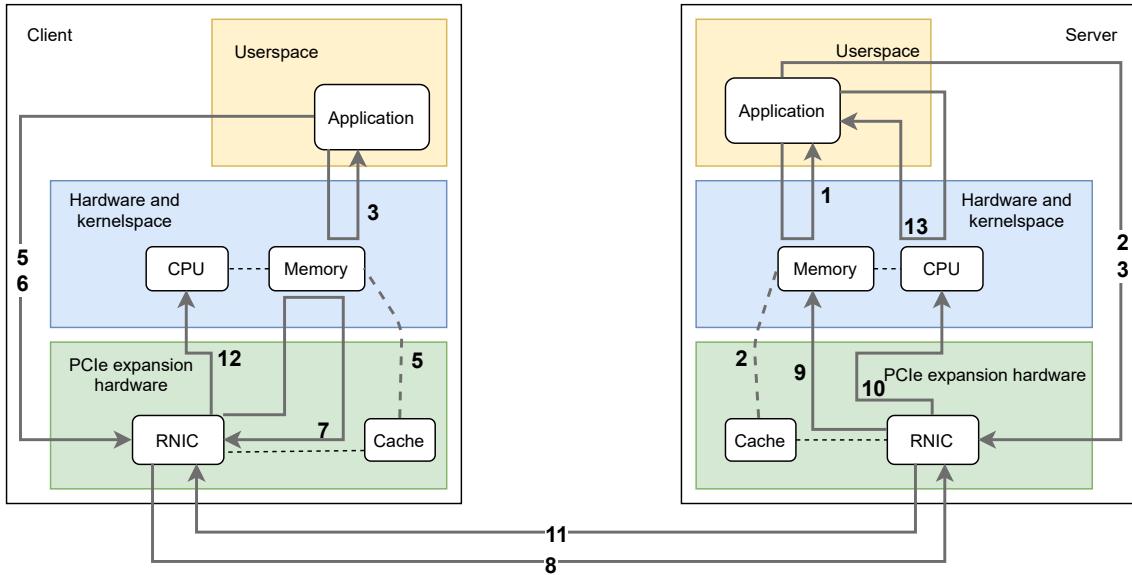


Figure 2.1: Process overview of pre-posting receive and send from client to server. Numbered arrows have a corresponding function described in 2.3.5. Numbered dashed lines are cache page entries for registered memory.

Unconnected datagrams also require metadata exchange. Data such as QP number, and physical port, which is needed to create AH, or join a multicast group. This can be achieved in the same way as with connected protocols.

2.3.5 Programming with RDMA

Unlike traditional sockets, much of RDMA’s interfacing is open to the programmer, allowing for detailed optimizations. This interface is developed and supported by the OpenFabrics Alliance[2]. The cost of this is the need for more memory management in userspace, which otherwise would be done by the kernel. In this subsection, this open interface will be discussed, what functions are needed to be done to make use of RDMA, and what happens internally. For this, figure 2.1 illustrates the process of an RC QP connection, and *SEND* and *RECV* operation, from client to server. It is assumed that client and server have successfully made a protection domain (PD), QP, and connect or exchanged data for these QP’s, see section 2.3.4 on how this can be done. The numbered arrows in figure 2.1 correspond to the following:

1. The server application should allocate memory for its receiving buffer.
2. This buffer should be registered in the RNIC under the PD used for this application.

2. BACKGROUND

By registering memory, a page table entry will be held in RNIC's cache, as shown with the corresponding dashed line.

3. Server pre-posts a receive. This is needed to be able to receive before the client will send. Usually this is done just before connecting, this way even if the receiving end lags behind, the RNIC will be ready to receive.
4. Client allocates the buffer that will be sent.
5. Also registering this buffer. This can be done while the server is completing steps 1 and 2.
6. Posting a *SEND* WR to the RNIC. This signals the RNIC that a given registered memory address will be sent via its QP.
7. The RNIC fetches this data in memory using DMA.
8. In step 8, this data is then sent over the network to the server.
9. The server's RNIC receives this data and performs an DMA operation to the buffer given by its *RECV* WR.
10. RNIC notifies with a completion event to the CPU, that it has a new event in its CQ.
11. RNIC also sends an completion event to the client.
12. This is passed along to the CPU.
13. The server can poll on the CQ, this tells the application that the *RECV* request has been completed, and can now be read with the same data as was in the clients buffer.

For more insight in which functions are available, see RDMA programming guide[26] or the OpenFabric Alliance[2] which develop kernel bypassing API for RDMA aware networks.

3

Designing a RDMA Key Value store

In this section, the design of the used key value store is shown, along with design decision made for RDMA. Further, benchmarking strategy will be explained, used to evaluate the performance of various RDMA transport types. Lastly, an overview of DAS-5 is given, which will be used to run benchmarks. This can be used when comparisons are made with other findings.

3.1 Key value store

This thesis is focused on the scalability using RDMA, and will not focus on advancing KV-stores. Therefore, a trivial KV store has been used. For this KV store, *SET* and *GET* instructions are mainly used, along with other commands for testing purposes. All possible commands can be seen in figure 3.1(b).

3.1.1 Requests

For a client to interact with a KV server, the client has to send a request towards the server. A request is structured as shown in figure 3.1(a). This key is used to identify the correct index within the hash table. A field for payload is always sent, however only used for *SET* commands. This is to have a constant data size, key having 64 bytes and message 256.

3.1.2 Response

After processing a request the server will always send a response. The structure of which is shown in 3.2(a). In case of a successful execution, an *OK* code will be given. For *GET* requests a payload is given along with this *OK* code, this is the value requested. Upon

3. DESIGNING A RDMA KEY VALUE STORE

Request	
int	client_id
method	method
key[64]	key
Msg[256]	message
size_t	key_len
size_t	msg_len
int	connection_closed

(a) Request structure

(b) Values for method

Figure 3.1: Structure for request with accommodated method

Response		response_code
response_code	code	OK
msg[256]	message	KEY_ERROR
size_t	msg_len	PARSING_ERROR
		STORE_ERROR
		SETOPT_ERROR
		UNK_ERROR

(a) Response structure

(b) Respond codes

Figure 3.2: Structure for response with accommodated response codes

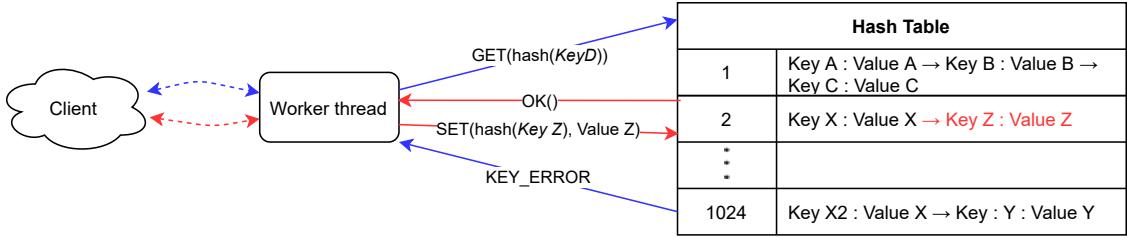


Figure 3.3: Key value server layout. Process worker thread takes with incoming requests, and what response is sent. A successful *SET* request is shown in red, adding a value "Value Z" to the internal linked list, with key "Key Z". Blue is an unsuccessful *GET* where the key is not known.

error, an error code will be sent back: *PARSING_ERROR*, or *UNKNOWN (UNK)*, latter of which is when the given command is unknown. Other response codes can be seen in 3.2(b)

3.1.3 Hash table

Internally, a hash table with 1024 buckets is used. Each bucket contains a linked list of key-value pairs. Figure 3.3 shows the structure of the KV store used. It should be noted, a linked list approach, as used here, does not scale well. As buckets are filled, either due to a hash collisions or extensive use, *GET* requests can cause significant delay, especially when encountering a key miss. This can be partially solved with a large hash table, this however will not be a long term solution, along with strong hash function. However, with a 95% *GET* request workload, the scalability issues will not show a significant performance loss. Moreover, in section 3.3 this issue will be kept constant throughout experiments.

3.1.4 Multithreaded

The KV server implemented for this thesis is also multithreaded. For every client a worker thread process will be created. This thread will read requests, process accordingly, and send back a response, all the while the main thread is used to accept and set up for new clients. In this way, worker threads are unaffected by new clients that are connecting.

To ensure concurrent and correct operation, each bucket has a mutex lock, and each key-value pair a read-write lock. This could cause performance loss as been seen in FaSST [24, 31].

3. DESIGNING A RDMA KEY VALUE STORE

3.2 RDMA

For this research, the two-sided verbs *SEND* and *RECV* are used for RDMA. Looking at table 2.1, across all transportation types *SEND/RECV* is the only available verb. This might come at a cost of performance, due to the channel schematic and minor CPU involvement. It has been shown that in some use cases, one-sided *WRITE* operations performs superior compared to *SEND*[22]. However, this difference is minimal, and the use of *WRITE* is not possible for this research.

To establish a baseline improvement and scalability, no optimizations are implemented. In section 6 possible optimizations are discussed.

3.2.1 Queue pairs

Connection based QP's will be handled similarly as in section 3.1.4 above. Every worker thread will have a client connected QP, which can only communicate with this client. This differs for UD. As stated in section 2.3, any UD QP can communicate with any other UD QP. This meaning, that the server only needs one QP for all clients. A worker thread is still created for every client, however, in the case of UD, all threads use a shared queue.

3.3 Benchmark design

To evaluate the performance and scalability of the RDMA KV store, a mult-client benchmark has been made. In total, 10 million tasks will be divided among clients. This number is constant to lessen the scalability issues with the underlying hash table. A task involves two operations, sending a request and receiving response. Only once a response is received from the server, can the client continue to the next task. This meaning that the full potential of RDMA is not being used, however it has been chosen for correctness, and such that latency can be evaluated accurately on a per-task level.

During client benchmark execution, the time will be taken at two points: before sending request, and after receiving response. With this, the throughput and latency of every client and every operation can be traced back. Time will be gathered with the *gettimeofday* function, and as accuracy achieves measurements at μ sec scale. These times will be kept in an array, and be returned after completing its tasks. The time will be later written to CSV files perform statistical analysis and graphs are drawn, as can be seen later on in the evaluation section 4. For this, Python 3.6[6] is used, along with pandas[5] and matplotlib[4].

3.3 Benchmark design

Cluster	Nodes	CPU type	Frequency (GHz)	Memory (GB)	Network
VU	68	dual 8-core	2.4	64	IB and GbE
LU	24	dual 8-core	2.4	64	IB and GbE
UvA	18	dual 8-core	2.4	64	IB and GbE
TUD	48	dual 8-core	2.4	64	IB and GbE
UvA-MN	31	dual 8-core	2.4	64	IB and GbE
ASTRON	9	dual 8/10/14-core	2.6	128/512	IB, 40 GbE, GbE

Table 3.1: DAS-5 cluster specifications

The benchmark is designed to evaluate each transportation types under similar conditions, and will follow the same path. Once a client is setup and connected with the server, it will start performing its tasks. These tasks follow the 30:1 *GET/SET* ratio that found by Atikoglu et. al.[10]. A *SET* request has a 5% chance of being generated. This is done by generating a random number as follows: $rand() \bmod 100 \leq 5$. Else a *GET* request is generated. The client sends out the request, and waits for response.

3.3.1 Experimental Setup

All performance tests and results have been gathered on the DAS-5 computing cluster[3]. This distributed system of computers spread across the Netherlands, and is used by research groups from VU Amsterdam, TU Delft, Leiden University, and many more. Each cluster varies slightly in specifications, however each, is equipped with dual Intel E5-2630v3 8-core CPUs, a 56 Gbit/s Inifiniband (IB) RDMA networking, and 1 Gbit/s classical ethernet networking. The specifications of each cluster is as shown in table 3.1.

All experiments make use of the IB network card, and is also configured to run TCP/IP. Furthermore, at least two nodes are used, his ensures that the server and clients are separated.

3. DESIGNING A RDMA KEY VALUE STORE

4

Evaluation

In this section, the experimental results are presented. The RDMA transportation types are compared within and against the baseline TCP implementation. Throughput and latency are mainly used to analysis the multicore scalability of these transportation types; with increasing number of clients, and therefore threads. For latency three figures have been used to evaluate: average latency, boxplot without outliers, and cumulative distribution function (CDF). The same raw data is used for both measurements, thus can be compared alongside each other.

The two designs are analyzed, with and without blocking for completion event. For TCP the same data has been used, as this is used as baseline.

In short, the results show:

- Without blocking, throughput reaches an peak, with all transportation types, at roughly 32 clients. UC performs best, with a maximum throughput of 370185 ops/sec at 33 clients. Further details are given in section 4.1.
- With blocking, throughput reaches an equilibrium up to 70 clients. UC, again, performance best, stabilizing at roughly x ops/sec.
- Latency increases across all transportation types, with increased number of clients, and both with and without blocking. Results are inversely comparable to throughput, again with UC performing best. Without blocking: UC has a average latency of roughly 49.2 μ sec at 32 clients and 179.7 μ sec at 60 clients. The spread also increases: standard deviation ranges from 46.6 μ sec at 32 clients, to 1628.2 μ sec at 60 clients. UC also performance best due to the stable 95th percentile, which is 110 μ sec for both 32 and 60 clients.

4. EVALUATION

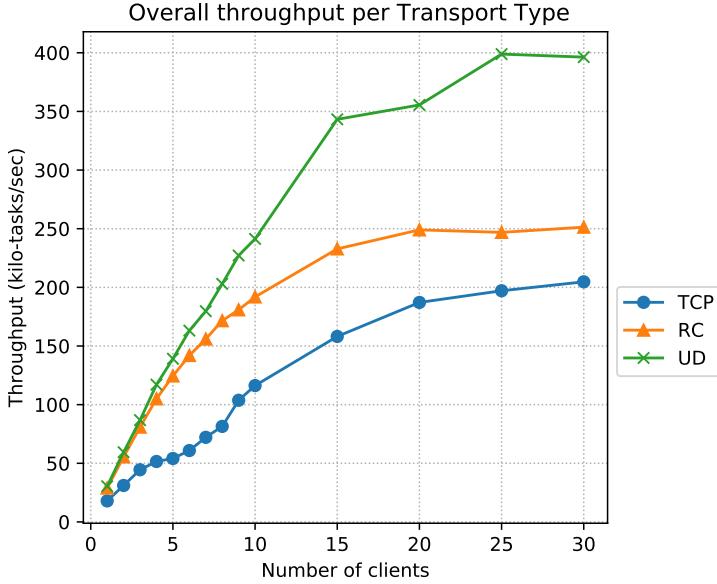


Figure 4.1: Throughput of clients executing 10 million operations

- With blocking: UD has a average latency of roughly $x \mu\text{sec}$ at 32 clients and $x \mu\text{sec}$ at 60 clients. The spread also increases: inner quartile range ranges from $x \mu\text{sec}$ at y clients, to $x \mu\text{sec}$ at y clients.

Section 4.2 delves deeper.

To recall the benchmarking setup: Ten million KV-store operations are divided equally between n -clients.

4.1 Throughput analysis

To compare the overall throughput between TCP, RC, UC, and UD with up to 70 clients. With non-blocking design, peak throughput performance is realised at roughly 32 clients. This can be seen in figure 4.1. For RC, maximum throughput is roughly 247782 ops/sec at 32 clients. While for UC, maximum throughput is roughly 370185 ops/sec at 33 clients, and for UD this is at roughly 357019 ops/sec at 31 clients. The maximum throughput TCP achieves in this number of clients range, is 206750 ops/sec at 33 clients. For RC, UC, and UD this is a percentage increase over TCP of: 19.85%, 79.05%, and 72.68%. However, after 32 clients, throughput has a sharp decline, with RC failing to connect properly above 60 clients, resulting in incomplete data at these points.

4.2 Latency analysis

Clients	TCP	RC	UC	UD
1				
2	6531.571861	24686.53216	24638.92191	26502.60875
3	13594.18984	23770.88044	24785.06225	28301.22973
4	4297.07597	24568.10076	27795.82235	23489.22267
5	7363.59922	17665.61705	19850.77766	22090.10711
6	6521.771444	16064.3204	20579.02917	24167.81283
7	11672.14625	13975.27874	21439.07072	16246.34785
8	18989.11049	14401.88884	17510.50509	14201.32075
9	12033.20216	4957.154159	17930.69447	21432.34552
10	11858.04359	18916.31359	20129.9112	17827.70575
15	9603.176325	8252.113695	19962.02981	16410.46889

Table 4.1: Throughput increase per additional client

4.1.1 Scalability rate

Multicore scalability of all RDMA transportation types, without blocking, is poor. Up to roughly 15 clients, throughput performance increases at a linear rate. For TCP, there is an average increase of 10246 ops/sec per additional client. For the RDMA transport types, for RC this is 16726 ops/sec per additional client, 21462 ops/sec for UC, and 21067 ops/sec for UD. This shows that unreliable transport types, UC and UD, scale better per additional client and worker thread, compared to RC and TCP. However, after 15 threads, performance begins to plateau for all transportation types and TCP, showing that the used KV-store is CPU bound. This is due to the pthread locking, used to provide isolation with concurrent threads operating on the KV-store.

After 32 clients and worker threads, multicore performance weakens, and dropping in performance. This is due to the non-blocked design, and 32 physical threads being the maximum on a single DAS-5 node. Without this blocking or scheduling, threads take CPU cycles while polling on the receive queue. With only a limited number of CPU cycles, this is not used efficiently, resulting in drop in performance for all worker threads.

4.2 Latency analysis

An important factor to KV-store is their quick response time. Latency should ideally scale similarly, preferably better, than throughput, such that with an increased number of clients, latency changes are unnoticed to the clients. RDMA offers fast response time

4. EVALUATION

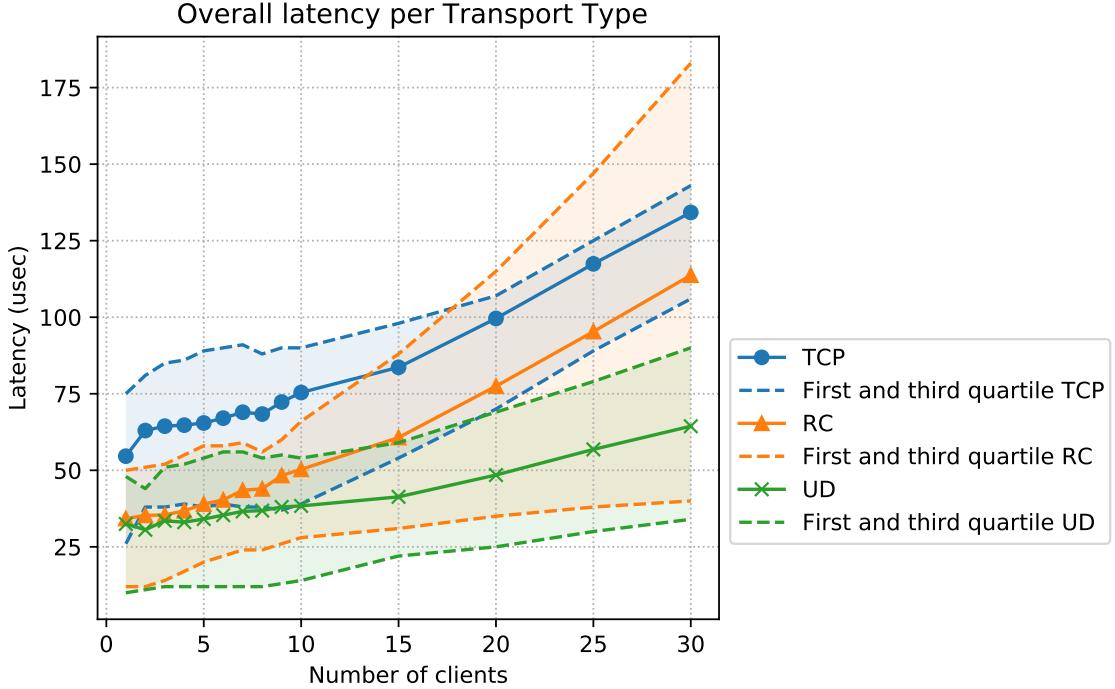


Figure 4.2: Latency of clients executing 10 million operations

due to the lack of CPU and kernel involvement. This can be seen in 4.3(b), where RDMA transportation types perform better than TCP up to 55 clients, after which UD performs worse, and RC following this trend.

At roughly 32 clients, relative performance is similar to throughput as seen in figure 4.1. Results show that TCP has an average latency of $138.0 \mu\text{sec}$ at 32 clients. RC has an average latency of $115.1 \mu\text{sec}$, a 16.61% improvement relative to TCP. UC performs best at 32 clients, with an average latency of $49.24 \mu\text{sec}$, which is an 64.33% improvement over TCP. Finally, UD has an average latency of $65.5 \mu\text{sec}$, which is an improvement of 52.5% over TCP.

It can also be observed that latency scales proportional with the number of clients for TCP, while this is only the case for RDMA up to roughly 32 clients. The average rate of change for TCP, RC, UC, and UD up to 32 clients is: $2.31 \mu\text{sec}$, $2.00 \mu\text{sec}$, $0.37 \mu\text{sec}$, and $0.72 \mu\text{sec}$. This shows that UC scales favorably up to 32 clients, closely followed by UD. After 32 clients, RDMA transportation types scale worse, with UD surpassing TCP's average latency at 60 clients. RC is trending towards surpassing TCP after 60 clients, however, this data could not be collected reliably. Lastly, UC still performs favorably compared to TCP at 70 clients, however, relative performance is now at 21.98%.

The latency between protocols is an inverse reflection of the throughput graph, scaling and overall performance. Similarly, TCP performed worst, with an average latency of 134.2 μ sec at 30 clients. RC hits an average of 113.7 μ sec, also at 30 clients. UD again performing best, reaching 64.4 μ sec on average.

Single client performance is lacking with HERD[22] and FaSST[24], and from what should be theoretically possible, being in the order of sub-10 μ sec latency. This could be due to lock contention with shared QP queues, in the case of UD, and/or poor performance of KV store. A further examination of this could support this.

Similarly to throughput, UD scales relatively well, while RC and TCP scale linearly. A near uniform latency for the first 15 clients, afterwards a gradual slope, compared to TCP and RC.

4.2.1 Variation from average

Consistent latency is important for clients, such that performance is predictable. Therefore, a box plot and CDF graph is used to examine the spread of latency at three points: 5 clients, 32 clients and 60 clients. Additionally, the median, inner-quartile range, 95th and 99th percentile, and standard deviation, is used to examine the spread numerically. This numeric data can be found in table 4.2. The complete table, with number of clients up to 70, can be found in table INSERT TABLE in the appendix.

From this it can be observed that from these three points, standard deviation increases as number of clients increase. With RDMA, there is a significant increase in variation after 32 clients. The median latency and 95th percentile of UC and UD remains similar between 32 and 60 clients, however the outliers or 99th percentile, increase significantly. This is further supported by the CDF figure of 32 clients INSERT FIGURE REF and 60 clients INSERT FIGURE REF. In these figures, 80% of the data is near 100 μ sec and 90 μ sec, for UD and UC respectively, deviating slightly between 32 and 60 clients. However, the top 1% of the data is not below the 500 μ sec shown for 60 clients, while this is the case for 32 clients. These outliers are caused by polling and concurrent number of threads, which exceed the number of physical threads. With high CPU usage, the time until a new request is seen and can be processed, increases, causing a delay for the client. These outliers can be also seen to cause the loss in throughput performance.

It can also be observed that RC suffers from a peak spread around 32 clients, which afterwards falls back down. Inner quartile range reaches a peak of 149 μ sec at 32 clients, with a median of 112 μ sec. It is unclear what causes this, however should be considered when reaching maximum number of physical threads.

4. EVALUATION

Clients	Median	IQR	95th percentile	99th percentile	Standard deviation
5	62	55	126	152	33.50
32	129	40	192	233	268.86
60	238	65	318	368	275.85

(a) TCP

Clients	Median	IQR	95th percentile	99th percentile	Standard deviation
5	31	43	88	99	29.33
32	112	149	227	231	83.10
60	48	56	120	7852.41	1662.83

(b) RC

Clients	Median	IQR	95th percentile	99th percentile	Standard deviation
5	29	49	84	95	29.26
32	39	70	110	122	46.59
60	35	56	110	1296	1629.16

(c) UC

Clients	Median	IQR	95th percentile	99th percentile	Standard deviation
5	29	46	84	95	30.45
32	57	64	126	146	65.59
60	54	65	143	9036	1693.72

(d) UD

Table 4.2: Numeric statistic latency. All statistical values have unit μ sec.

4.2.2 Distribution of average latency

Looking at the box plots depicted in INSERT FIGURE and CDF graphs in INSERT FIGURES, TCP is normally distributed at 32 and 60 clients, while UC and UD are skewed left. This can be seen in the sharp increase in cumulative probability in the CDF graphs INSERT FIGURES. The distribution of RC latency results varies, as it is normally distributed at skewed left at 5 and 60 clients.

In figure 4.2 the variation is given in form of first and third quartile. It can be seen that TCP and UD have a constant, a small, variation. TCP is skewed towards higher latencies at 30 clients, while UD, and RC, stay normally distributed throughout. RC has a rapidly increasing variation, which is concerning for scalability and consistency. As number of clients increases, the realised latency could vary drastically between tasks. These issues could possibly be due to context switching, more in subsection 4.2.4 below.

4.2.3 All 30 clients and per operation latency

Figure 4.3 delves deeper into the details of the latency outliers that are present. Note on the scale, not similar between protocols, and compared to figure 4.2. These outliers are the extreme outliers, not displayed in figure 4.3(b). TCP and RC see spikes up to $114185 \mu\text{sec}$, or 114.2ms , and $32046 \mu\text{sec}$, or 32.05ms . UD on the other only spikes to a maximum of $11702 \mu\text{sec}$, or 11.7ms . These outliers have a significant impact on the realised performance. In the case of TCP and UD, these seem to occur at random. For RC this does not seem to be the case, which suggests a systematic error is at cause.

4.2.4 RDMA multiple QP's causing cache misses and context switching

In section 2.3.2 the issue with holding multiple QP's within the RNIC, has been discussed. The RNIC needs to perform context switching when changing QP. As the number of QP's increase, the number of context switches needed increases. Along this, the chance of cache misses are also increased. This decreases the performance, with increased latency, and in throughput as shown here. With context switching, a large delay occurs, and harms the performance per client, but also of the overall network. For this reason, for large number of clients, a one-to-one QP, such as RC and UC, is not advisable.

4.2.4.1 Cache misses as increased variation for RC

Furthermore, cache misses occur more often at the starting phase of KV store operations, as can be seen in section 4.3(b) above. RC has consistently higher latency (0.75ms) for

4. EVALUATION

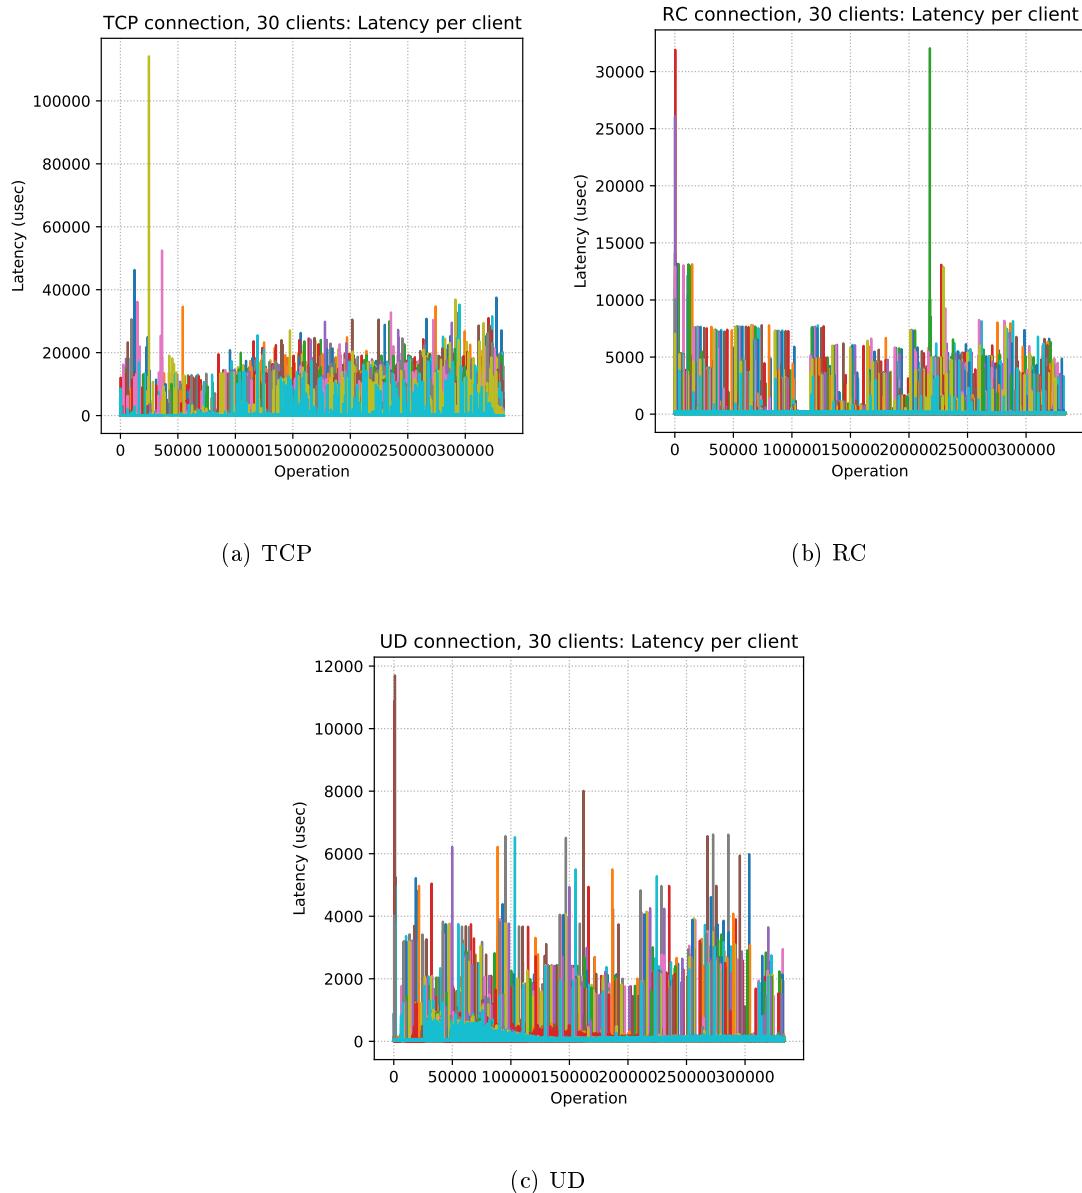


Figure 4.3: Structure for response with accommodated response codes

4.2 Latency analysis

some clients at the start. Initially it was thought to be retransmissions, however for this was set to 67.1ms. Since this only occurs to some clients, and not all, this could be a result from the aforementioned context switching and RNIC cache misses. At the start of operations, worker threads are operating in sync. This would result in RNIC processing multiple WR's concurrently. Some QP's will be cached and experience little cache misses throughout the starting phase, while others would encounter cache misses repeatedly. Once threads are out of sync and KV store operation latency increases (due to the increasing size of the hash table), the cache misses will occur for all QP's. This can be seen in figure 4.3(b), as the latency in the final phase is randomized and less constant, compared to the starting phase.

With cache misses, latency per client varies, as some QP's would require further lookup after cache miss. This effect can be seen in figure 4.2, as increasing number of clients, and therefore QP's, increases the variation dramatically. Further investigation into cache misses would support these claims, this has not been done for this thesis.

4. EVALUATION

5

Related Work

5.1 HERD

There are several proposed RDMA KV store designs. One of which, HERD [22], uses a combination of transport types. HERD uses RDMA *WRITE* over UC for requests and *SEND* over UD for responses. Kalia et. al. have shown that incoming *WRITES* offers lower latency and higher throughput compared to *READ*. Outgoing *WRITES* have also been shown to not scale well, making this unadvisable for sending responses.

HERD out performs the RDMA KV store presented in this thesis significantly. This could due to multiple factors:

- HERD has implemented some optimizations towards prefetching before posting a *SEND WR*. With this they have observed roughly 30% improvement, in the best case comparison (2 random memory accesses at 4 CPU cores.)
- Making use of one-sided *WRITE* for requests, bypassing the CPU.
- Using inlined data for key. Inlining small payloads decreases latency as there is no need for a DMA operation.

However, Kalia et. al. stated that performance should be comparable to *SEND/SEND*, given that inlining is possible. It was found that for a large number of clients and/or requests, the round-robin polling used, is inefficient. With 1000s of clients, a *SEND/SEND* design would scale better than the *WRITE/SEND* used in HERD.

5. RELATED WORK

6

Future improvements

6.1 Experimental transport types

Current transportation types all have advantages and disadvantages. Reliable transport has advantages being predictable packet behavior. However, as has been shown in this thesis, connection transports (such as RC) have scalability issues, due to the RNIC caching issues with increasing number of QP's. Dynamically Connected Transport (DCT) is, currently, an experimental transport type which addresses these issues. DCT has a connection based design, while only requiring one QP. This is achieved by dynamically connecting and disconnecting with remote QP's, which would introduce additional latency with short-lived clients, or when dealing with concurrent large number of clients. However, could have the potential to combine the reliability of RC and

6.2 Optimizations

Currently, no optimizations are used to improve performance. Since this thesis focused on performance across transport types, this needed to remain consistent throughout. Optimizations, such as those used in HERD, making use of prefetching and inlined data, improve performance significantly. Other optimizations can be found in Kalia et. al. paper on design guidelines and possible optimizations[23]. These are hardware focused, including batching, reducing the number of PIO and DMA operations.

This thesis has shown the potential of using UD as scalable and high throughput transport type for RDMA KV stores, however this could be improved further. One improvement is to have multiple groups of clients, each with a single QP. This has been shown to be effective in ScaleRPC[11]. Grouping clients reduces RNIC cache contention, by limiting

6. FUTURE IMPROVEMENTS

the number of worker threads competing for cache. However, like with RC and UC, using multiple QP's increases context switching, which harms performance. ScaleRPC have proposed solutions to this by "warming up" QP's.

6.3 Newer hardware

Performance benchmarks for this thesis have been ran on DAS-5. DAS-5 is becoming outdated, with the newer DAS-6 being available mid-2021[8]. DAS-6 will make use of 100 Gbit/s Infiniband RNIC, near 2x bandwidth compared to DAS-5.

Conclusion

RDMA has been shown to be a promising advancement for key values stores. The low latency that RDMA brings with it, opens up new possibilities and improved performance. RDMA also requires more design decision to be made, which transportation protocol to use, which optimizations to use, and which verbs to use. In this thesis, the transportation protocols RC and UD have been compared, on performance, to TCP.

RQ1: In industry, KV stores require to be scalable and low latency. Connection based transportation (RC and UC) put stress on RNIC when using large number of clients, hindering performance with respect to throughput and scalability. RC has similar scalability to TCP, with a near 60 kilo-tasks/sec increase in throughput across all number of clients. Best performing transportation type is UD, with close to 400 kilo-tasks/sec throughput and 64 μ sec at 30 clients. It has been shown that UD offers a 58.7% improvement in throughput against RC, and 45.5% improvement in latency. For scalability UD is recommended to use, as this offers the best overall performance and sustains this up to 30 clients, and has the potential for optimizations to further improve performance.

RQ2: The advantages and disadvantages present with the different transportation types impact the KV store design choices. In context of KV stores, UD is a compelling transportation type, however does not offer RDMA verbs such as *READ* and *WRITE*, which could offer for improved performance. UD also allow for more optimizations, due to their one-to-many QP. UD performance in this thesis have found to be limited. Multiple UD QP's along with client grouping could continue scalability further. Connection based transportation types are limited by one-to-one QP, although also have some room for improvements.

7. CONCLUSION

Appendix

7. CONCLUSION

References

- [1] memcached: free and open source, high-performance, distributed memory object caching system. <http://memcached.org>, 2003. Online: accessed 28 May 2021. 5
- [2] Openfabrics alliance. <https://www.openfabrics.org/>, 2004. OpenFabrics Alliance: develops, tests, licenses, supports and distributes RDMA/Advanced Networks software. 9, 10
- [3] Das-5: Distributed asci supercomputer. <https://www.cs.vu.nl/das5/home.shtml>, 2012. Online: accessed 28 May 2021. 15
- [4] Matplotlib: Visualization with python. <https://matplotlib.org/>, 2016. Python graphing and visualization tool. 14
- [5] Pandas. <https://pandas.pydata.org/>, 2016. Data analysis and manipulation tool. 14
- [6] Python. <https://www.python.org/>, 2016. Python programming language. 14
- [7] Rdma exmaple. <https://github.com/animeshtrivedi/rdma-example>, 2019. RDMA example implementation used as inspiration for RC implementation and RDMA helper function. 4
- [8] Das-6: Distributed asci supercomputer. <https://www.cs.vu.nl/das6/home.shtml>, 2021. Online: accessed 28 May 2021. 30
- [9] Rdma based key-value store. <https://github.com/jtro0/RDMA-KV-store>, 2021. All relevant thesis project files are openly available as of June 2021. 3
- [10] Berk Atikoglu, Yuehai Xu, Eitan Frachtenberg, Song Jiang, and Mike Paleczny. Workload analysis of a large-scale key-value store. In *Proceedings of the 12th ACM SIGMETRICS/PERFORMANCE joint international conference on Measurement and Modeling of Computer Systems*, pages 53–64, 2012. 3, 5, 15

REFERENCES

- [11] Youmin Chen, Youyou Lu, and Jiwu Shu. Scalable rdma rpc on reliable connection with efficient resource sharing. In *Proceedings of the Fourteenth EuroSys Conference 2019*, pages 1–14, 2019. 1, 2, 29
- [12] Giuseppe DeCandia, Deniz Hastorun, Madan Jampani, Gunavardhan Kakulapati, Avinash Lakshman, Alex Pilchin, Swaminathan Sivasubramanian, Peter Vosshall, and Werner Vogels. Dynamo: Amazon’s highly available key-value store. *ACM SIGOPS operating systems review*, 41(6):205–220, 2007. 1, 2
- [13] Aleksandar Dragojević, Dushyanth Narayanan, Miguel Castro, and Orion Hodson. Farm: Fast remote memory. In *11th USENIX Symposium on Networked Systems Design and Implementation (NSDI 14)*, pages 401–414, Seattle, WA, April 2014. USENIX Association. ISBN 978-1-931971-09-6. URL <https://www.usenix.org/conference/nsdi14/technical-sessions/dragojević>. 1
- [14] Robert Escriva, Bernard Wong, and Emin Gün Sirer. Hyperdex: A distributed, searchable key-value store. In *Proceedings of the ACM SIGCOMM 2012 conference on Applications, technologies, architectures, and protocols for computer communication*, pages 25–36, 2012. 1, 5
- [15] Philip Werner Frey and Gustavo Alonso. Minimizing the hidden cost of rdma. In *2009 29th IEEE International Conference on Distributed Computing Systems*, pages 553–560. IEEE, 2009. 6, 8
- [16] Roxana Geambasu, Amit A Levy, Tadayoshi Kohno, Arvind Krishnamurthy, and Henry M Levy. Comet: An active distributed key-value store. In *OSDI*, pages 323–336, 2010. 1
- [17] Richard W Hamming and Roger Pinkham. The art of doing science and engineering: Learning to learn. *Mathematical Intelligencer*, 20(3):60, 1998. 3
- [18] Nathan Hanford, Vishal Ahuja, Matthew K Farrens, Brian Tierney, and Dipak Ghosal. A survey of end-system optimizations for high-speed networks. *ACM Computing Surveys (CSUR)*, 51(3):1–36, 2018. 5, 6
- [19] Gernot Heiser. Systems benchmarking crimes. <https://www.cse.unsw.edu.au/~gernot/benchmarking-crimes.html>, 2010. "Online: accessed 28 May 2021". 3

REFERENCES

- [20] Alexandru Iosup, Laurens Versluis, Animesh Trivedi, Erwin Van Eyk, Lucian Toader, Vincent Van Beek, Giulia Frascaria, Ahmed Musaafir, and Sacheendra Talluri. The atlare vision on the design of distributed systems and ecosystems. In *2019 IEEE 39th International Conference on Distributed Computing Systems (ICDCS)*, pages 1765–1776. IEEE, 2019. 3
- [21] Raj Jain. *The art of computer systems performance analysis: techniques for experimental design, measurement, simulation, and modeling.* John Wiley & Sons, 1990. 3
- [22] Anuj Kalia, Michael Kaminsky, and David G Andersen. Using rdma efficiently for key-value services. In *Proceedings of the 2014 ACM Conference on SIGCOMM*, pages 295–306, 2014. 1, 2, 7, 14, 21, 27
- [23] Anuj Kalia, Michael Kaminsky, and David G Andersen. Design guidelines for high performance {RDMA} systems. In *2016 {USENIX} Annual Technical Conference ({USENIX} {ATC} 16)*, pages 437–450, 2016. 1, 2, 6, 29
- [24] Anuj Kalia, Michael Kaminsky, and David G Andersen. Fasst: Fast, scalable and simple distributed transactions with two-sided ({RDMA}) datagram rpcs. In *12th {USENIX} Symposium on Operating Systems Design and Implementation ({OSDI} 16)*, pages 185–201, 2016. 2, 7, 13, 21
- [25] Hyeontaek Lim, Dongsu Han, David G Andersen, and Michael Kaminsky. {MICA}: A holistic approach to fast in-memory key-value storage. In *11th {USENIX} Symposium on Networked Systems Design and Implementation ({NSDI} 14)*, pages 429–444, 2014. 1, 5
- [26] *RDMA Aware Networks Programming User Manual.* Mellanox Technologies, 350 Oakmead Parkway Suite 100, Sunnyvale, CA, USA, 1.7 edition, 2015. Online: accessed 28 May 2021. 10
- [27] Christopher Mitchell, Yifeng Geng, and Jinyang Li. Using one-sided {RDMA} reads to build a fast, cpu-efficient key-value store. In *2013 {USENIX} Annual Technical Conference ({USENIX} {ATC} 13)*, pages 103–114, 2013. 2
- [28] John Ousterhout. Always measure one level deeper. *Communications of the ACM*, 61(7):74–83, 2018. 3

REFERENCES

- [29] John Ousterhout, Parag Agrawal, David Erickson, Christos Kozyrakis, Jacob Lev-erich, David Mazières, Subhasish Mitra, Aravind Narayanan, Guru Parulkar, Mendel Rosenblum, et al. The case for ramclouds: scalable high-performance storage entirely in dram. *ACM SIGOPS Operating Systems Review*, 43(4):92–105, 2010. 5
- [30] Ken Peffers, Tuure Tuunanen, Marcus A Rothenberger, and Samir Chatterjee. A design science research methodology for information systems research. *Journal of management information systems*, 24(3):45–77, 2007. 3
- [31] Haonan Qiu, Xiaoliang Wang, Tianchen Jin, Zhuzhong Qian, Baoliu Ye, Bin Tang, Wenzhong Li, and Sanglu Lu. Toward effective and fair rdma resource sharing. In *Proceedings of the 2nd Asia-Pacific Workshop on Networking*, pages 8–14, 2018. 7, 13
- [32] Sameer Seth and M Ajaykumar Venkatesulu. *TCP/IP Architecture, Design, and Implementation in Linux*, volume 68. John Wiley & Sons, 2009. 5
- [33] Vijay Vasudevan, Amar Phanishayee, Hiral Shah, Elie Krevat, David G Andersen, Gregory R Ganger, Garth A Gibson, and Brian Mueller. Safe and effective fine-grained tcp retransmissions for datacenter communication. *ACM SIGCOMM computer communication review*, 39(4):303–314, 2009. 7
- [34] Kenichi Yasukata, Michio Honda, Douglas Santry, and Lars Eggert. Stackmap: Low-latency networking with the {OS} stack and dedicated nics. In *2016 {USENIX} Annual Technical Conference ({USENIX} {ATC} 16)*, pages 43–56, 2016. 6